

Can a Tiny AI on an 8 GB MacBook Write Real Business Logic?

Locating the spec-dialect boundary where a 3-billion-parameter coder model stops producing correct code

Muhammad Junaid — iteration, ladder experiments & repair loops · Muhammad Ansar — original research & problem · 2 July 2026

ABSTRACT

We ask whether Qwen2.5-Coder-3B-Instruct-4bit, run locally on Apple MLX, can generate a `process_calculations()` that scores a perfect 12/12 on a ten-rule pricing-and-commission specification. The model is fixed; the only lever is the *spec*. A re-dialected spec (v2) closes the gap from the original's 1/12 to **12/12 on the first greedy attempt**. We then build an eight-rung *spec ladder* — same ten rules and same targets, varying only the phrasing — to find the exact point of failure. The break is sharp and surprising: it is not the two constructs originally blamed (a multi-tier conditional and a truthy-string boolean), but the removal of the literal lookup tables, after which the 3B can no longer assemble the multi-step *commission* calculation. Three control rungs confirm it: keep the tables and the model tolerates every "trap." Finally, a white-box repair loop that feeds per-step intermediate diffs rescues the boolean rung to 12/12 at a low temperature, but does not reliably rescue the fully de-scaffolded rung — reaffirming that for a small model the dependable lever is the specification, not a cleverer repair loop.

1 The question and the setup

Small, locally-hosted language models are attractive: no data leaves the laptop, no per-token bill, and an 8 GB Apple Silicon MacBook can host a 3B model comfortably. The open question is whether such a model can produce *numerically correct* business logic — not plausible-looking code, but code whose output matches ground truth to the cent, on every row.

The task is fixed throughout. A model must emit a Python function `process_calculations(csv_path)` that applies a ten-rule pricing and commission ruleset to each row of a CSV and returns the running total `total_receivables`. Success is **12/12**: all twelve ground-truth rows matched within a 0.5 tolerance. The model is fixed at Qwen2.5-Coder-3B-Instruct-4bit (the *Coder* variant, via `mlx-lm`); we never swap in a larger or different model. Every experiment holds the model, the harness, and the arithmetic constant, and changes only the wording of the spec.

2 The initial win — and why the old solution failed

A five-rule version of the spec, each rule a single unambiguous operation, is solved by the 3B on the first try: **12/12, no repairs**. The trouble begins at ten rules. Running the article's original *True/False* spec, the model's best attempts are **3/12 on the 3B** and **1/12 on the 7B** — the architecture is right but every row is 2–5% off. Two failures, both *linguistic* rather than a capability ceiling, explain it:

- **Rule 4 — the volume tier.** The spec's three-way tier (`<100 → 0%`, `<1000 → 5%`, `≥1000 → 10%`) is collapsed by the model into `0.05 if v<100 else (0.1 if v<1000 else 0.1)` — wrong in two of three branches.

- **Rule 10 — the boolean.** The `tax verified` cell is the string `"False"`, which is *truthy* in Python, so `not row["tax verified"]` skips withholding at exactly the rows where it should apply.

The fix (the "v2" spec) re-expresses the *same ten rules* in the model's dialect: Rule 4 becomes two independent additive checks (`if v>=100: +0.05; if v>=1000: +0.05`) with an explicit "no elif"; the boolean becomes `yes / no` text in both spec and data; rates are handed over as literal Python dictionaries and bare decimals. The arithmetic is bit-identical to the original (maximum absolute difference in targets = 0.0). The result: **12/12 on the first greedy attempt**. Same problem, phrased so the model cannot trip — the lever is the spec, not the model.

3 The sweet-spot experiment

The v2 spec passes and the original fails, but those are two distant points. *Where exactly* is the boundary, and how technical must a spec author be to stay on the right side of it? To answer precisely we build a ladder of eight specs. Every rung describes the identical ten-rule problem against identical targets; the harness is frozen at the configuration that produced the validated 12/12 (same instruction wrapper, same decoding, same scorer). **Only the spec file changes**, so any change in outcome is attributable to the spec's dialect alone.

Rungs L1–L5 strip one layer of authoring scaffolding at a time, from the heavily-engineered v2 dialect down to the natural original. Rungs L6–L8 are *controls*: they keep the literal lookup tables but re-arm the volume-tier and boolean "traps," to test whether those traps are what actually break the model. Each level is scored by **best-of-ten** sampled attempts at temperature 0.4; a rung is "reachable" if any of the ten attempts reaches a full 12/12 — the meaningful capability threshold for a model this single-shot-noisy.

Figure 1 (next page) shows all eight specs as the model sees them, scaled to a single page. The prose after it walks through each rung.

Figure 1 · The eight-rung spec ladder (exactly the text fed to the model, HTML authoring notes stripped)

L1 · Prescriptive pseudocode <pre># Pricing & commission calculation spec Write `process_calculations(csv_path)` that reads a CSV and returns a list of dicts, one per row, each containing the key `total_receivables`. ## Input columns Each row has these columns (header names may have extra surrounding spaces (strip leading/trailing whitespace only, and keep inner spaces like in `payment` type); compare text in lower case; convert `volume` and `productid` to numbers; an empty cell means the stated default): - `productid` : "1", "2", or "3" - `status` : "individual" or "corporate" - `volume` : an integer - `season` : "spring", "summer", "winter", or "autumn" - `customer` : "strategic" or "non-strategic" - `region` : "emea" or "non-emea" - `payment type` : "ach/wire" or "credit-card" - `transaction type` : "domestic", "international", or "cross-border" - `product category` : "electronics", "groceries", or "luxury" - `tax verified` : "yes" or "no" (an empty cell counts as "no") ## Lookup tables (copy these exactly) PRODUCT_RATES = { 1: {"individual": 1000, "corporate": 900}, 2: {"individual": 1500, "corporate": 1200}, 3: {"individual": 2000, "corporate": 1700}, } CATEGORY_MULTIPLIER = {"electronics": 1.2, "groceries": 0.8, "luxury": 1.5} ## The 10 rules (apply in this exact order) 1. **Rate.** Look up `rate` = PRODUCT_RATES[productid][status]`. 2. **Base.** `receivable_before_discount` = volume * rate. 3. **Spring discount rate.** Start `spring_rate` = 0.0. If `season` is "spring", set `spring_rate` = 0.05`. 4. **Volume discount rate.** Build it up with two independent checks (do NOT use elif): - Start `volume_rate` = 0.0. - If `volume` >= 100, add `0.05` to `volume_rate`. - If `volume` >= 1000, add another `0.05` to `volume_rate`. - (So volume under 100-999 gives 0.05, 1000 or more gives 0.10.) 5. **Strategic discount rate.** Start `strategic_rate` = 0.0. If `customer` is "strategic", set `strategic_rate` = 0.05. 6. **Apply discounts.** `total_discount_rate` = `spring_rate` + `volume_rate` + `strategic_rate`. `receivable_after_discount` = `receivable_before_discount` * (1 - `total_discount_rate`). 7. **Region loading.** Start `region_loading` = 0.0. If `region` is "emea", set `region_loading` = 0.10 * `receivable_after_discount`. 8. **Payment loading.** Start `payment_loading` = 0.0. If `payment type` is "credit-card", set `payment_loading` = 0.025 * `receivable_after_discount`. 9. **Commission.** Look up `multiplier` = CATEGORY_MULTIPLIER[product category]`.   `commission` = (region_loading + payment_loading) * multiplier`. 10. **Tax withholding and total.** `subtotal_receivable_after_discount` + `commission`. Start `withholding` = 0.0. If `tax verified` is "no" (or the cell was empty), set `withholding` = 0.03 * `subtotal`. If `tax verified` is "yes", leave `withholding` = 0.0. `total_receivables` = `subtotal` - `withholding`. Output: a List of dicts, one per input row, each with the key `total_receivables`.</pre>	L5 · Original dialect (boolean) <pre># Pricing & commission calculation spec Write `process_calculations(csv_path)` that reads the CSV and returns one dict per row with the key `total_receivables`. ## Input columns Each row has these columns (header names may have extra surrounding spaces (strip leading/trailing whitespace only, and keep inner spaces like in `payment` type); compare text in lower case; convert `volume` and `productid` to numbers; an empty cell means the stated default): - `productid` : "1", "2", or "3" - `status` : "individual" or "corporate" - `volume` : an integer - `season` : "spring", "summer", "winter", or "autumn" - `customer` : "strategic" or "non-strategic" - `region` : "emea" or "non-emea" - `payment type` : "ach/wire" or "credit-card" - `transaction type` : "domestic", "international", or "cross-border" - `product category` : "electronics", "groceries", or "luxury" - `tax verified` : True or False (an empty cell counts as False) ## The calculation The per-unit price depends on `productid` and `status`. For productid 1 it is \$1,000 when status is individual and \$900 when corporate; for productid 2 it is \$1,500 / \$1,200; for productid 3 it is \$2,000 / \$1,700. The base receivable is `volume` times that per-unit price. Then three discounts come off, and they stack (add together): - **5%* if the `season` is spring; - a volume discount based on how much was ordered: under 100 units gets no volume discount, 100 to 999 units gets 5%, and 1000 units or more gets 10%; - **5%* if the `customer` is strategic. Apply the combined discount to get the receivable after discount: Two surcharges are then figured from the receivable after discount: - **0%* if the `region` is EMEA; - **2.5%* if the `payment type` is credit-card. Commission is the sum of those two surcharge amounts, scaled by a factor that depends on `product category`: electronics x1.2, groceries x0.8, luxury x1.5. The subtotal is the receivable after discount plus the commission. The `tax verified` column is True or False; when the sale has not been tax-verified, withhold 3% of the subtotal; when it has been verified, withhold nothing. `total_receivables` is the subtotal minus the withholding.</pre>
L2 · Declarative spec <pre># Pricing & commission calculation spec `process_calculations(csv_path)` reads a CSV and returns one dict per row, each with the key `total_receivables`. ## Input columns Each row has these columns (header names may have extra surrounding spaces (strip leading/trailing whitespace only, and keep inner spaces like in `payment` type); compare text in lower case; convert `volume` and `productid` to numbers; an empty cell means the stated default): - `productid` : "1", "2", or "3" - `status` : "individual" or "corporate" - `volume` : an integer - `season` : "spring", "summer", "winter", or "autumn" - `customer` : "strategic" or "non-strategic" - `region` : "emea" or "non-emea" - `payment type` : "ach/wire" or "credit-card" - `transaction type` : "domestic", "international", or "cross-border" - `product category` : "electronics", "groceries", or "luxury" - `tax verified` : "yes" or "no" (an empty cell counts as "no") ## Lookup tables PRODUCT_RATES = { 1: {"individual": 1000, "corporate": 900}, 2: {"individual": 1500, "corporate": 1200}, 3: {"individual": 2000, "corporate": 1700}, } CATEGORY_MULTIPLIER = {"electronics": 1.2, "groceries": 0.8, "luxury": 1.5} ... ## The calculation The base receivable is `volume * PRODUCT_RATES[productid][status]`. Three discount rates then apply, and they add together: - a spring rate of 0.05 when season is "spring", otherwise 0.0; - a volume rate that is 0.05 for volume >= 100 plus an additional 0.05 for volume >= 1000 (the two thresholds are independent and stack, so volume >= 1000 earns the full 0.10); - a strategic rate of 0.05 when `customer` is "strategic", otherwise 0.0. The receivable after discount is the base times (1 - sum of those three rates). Two loadings are computed from the receivable after discount: - a region loading of 0.10 * receivable_after_discount when `region` is "emea", else 0.0; - a payment loading of 0.025 * receivable_after_discount when `payment type` is "credit-card", else 0.0. The commission is (region_loading + payment_loading) * CATEGORY_MULTIPLIER[product category]. The subtotal is `receivable_after_discount` + `commission`. Tax withholding is 0.03 * subtotal when `tax verified` is "no" (or empty), and 0.0 when it is "yes". `total_receivables` = subtotal - withholding.</pre>	L6 · Tables + natural tier <pre># Pricing & commission calculation spec `process_calculations(csv_path)` reads a CSV and returns one dict per row, each with the key `total_receivables`. ## Input columns Each row has these columns (header names may have extra surrounding spaces (strip leading/trailing whitespace only, and keep inner spaces like in `payment` type); compare text in lower case; convert `volume` and `productid` to numbers; an empty cell means the stated default): - `productid` : "1", "2", or "3" - `status` : "individual" or "corporate" - `volume` : an integer - `season` : "spring", "summer", "winter", or "autumn" - `customer` : "strategic" or "non-strategic" - `region` : "emea" or "non-emea" - `payment type` : "ach/wire" or "credit-card" - `transaction type` : "domestic", "international", or "cross-border" - `product category` : "electronics", "groceries", or "luxury" - `tax verified` : "yes" or "no" (an empty cell counts as "no") ## Lookup tables PRODUCT_RATES = { 1: {"individual": 1000, "corporate": 900}, 2: {"individual": 1500, "corporate": 1200}, 3: {"individual": 2000, "corporate": 1700}, } CATEGORY_MULTIPLIER = {"electronics": 1.2, "groceries": 0.8, "luxury": 1.5} ## The calculation The base receivable is `volume * PRODUCT_RATES[productid][status]`. Three discount rates then apply, and they add together: - a spring rate of 0.05 when season is "spring", otherwise 0.0; - a volume rate based on how much was ordered: 0.0 under 100 units, 0.05 from 100 to 999 units, and 0.10 at 1000 units or more; - a strategic rate of 0.05 when `customer` is "strategic", otherwise 0.0. The receivable after discount is the base times (1 - sum of those three rates). Two loadings are computed from the receivable after discount: - a region loading of 0.10 * receivable_after_discount when `region` is "emea", else 0.0; - a payment loading of 0.025 * receivable_after_discount when `payment type` is "credit-card", else 0.0. The commission is (region_loading + payment_loading) * CATEGORY_MULTIPLIER[product category]. The subtotal is `receivable_after_discount` + `commission`. Tax withholding is 0.03 * subtotal when `tax verified` is "no" (or empty), and 0.0 when it is "yes". `total_receivables` = subtotal - withholding.</pre>
L3 · Business prose <pre># Pricing & commission calculation spec Write `process_calculations(csv_path)` that reads the CSV and returns one dict per row with the key `total_receivables`. ## Input columns Each row has these columns (header names may have extra surrounding spaces (strip leading/trailing whitespace only, and keep inner spaces like in `payment` type); compare text in lower case; convert `volume` and `productid` to numbers; an empty cell means the stated default): - `productid` : "1", "2", or "3" - `status` : "individual" or "corporate" - `volume` : an integer - `season` : "spring", "summer", "winter", or "autumn" - `customer` : "strategic" or "non-strategic" - `region` : "emea" or "non-emea" - `payment type` : "ach/wire" or "credit-card" - `transaction type` : "domestic", "international", or "cross-border" - `product category` : "electronics", "groceries", or "luxury" - `tax verified` : "yes" or "no" (an empty cell counts as "no") ## The calculation The per-unit price depends on `productid` and `status`. For productid 1 it is \$1,000 when status is individual and \$900 when corporate; for productid 2 it is \$1,500 / \$1,200; for productid 3 it is \$2,000 / \$1,700. The base receivable is `volume` times that per-unit price. Then three discounts come off, and they stack (add together): - **5%* if the `season` is spring; - a volume discount of **5%* once `volume` reaches 100, plus a further **5%* once `volume` reaches 1000 (so 100-999 is 5% and 1000-or-more is 10%); - **5%* if the `customer` is strategic. Apply the combined discount to get the receivable after discount: Two surcharges are then figured from the receivable after discount: - **0%* if the `region` is EMEA; - **2.5%* if the `payment type` is credit-card. Commission is the sum of those two surcharge amounts, scaled by a factor that depends on `product category`: electronics x1.2, groceries x0.8, luxury x1.5. The subtotal is the receivable after discount plus the commission. If `tax verified` is "no" (or blank), withhold **3%* of the subtotal; if it is "yes", withhold nothing. `total_receivables` is the subtotal minus the withholding.</pre>	L7 · Tables + True/False <pre># Pricing & commission calculation spec `process_calculations(csv_path)` reads a CSV and returns one dict per row, each with the key `total_receivables`. ## Input columns Each row has these columns (header names may have extra surrounding spaces (strip leading/trailing whitespace only, and keep inner spaces like in `payment` type); compare text in lower case; convert `volume` and `productid` to numbers; an empty cell means the stated default): - `productid` : "1", "2", or "3" - `status` : "individual" or "corporate" - `volume` : an integer - `season` : "spring", "summer", "winter", or "autumn" - `customer` : "strategic" or "non-strategic" - `region` : "emea" or "non-emea" - `payment type` : "ach/wire" or "credit-card" - `transaction type` : "domestic", "international", or "cross-border" - `product category` : "electronics", "groceries", or "luxury" - `tax verified` : True or False (an empty cell counts as False) ## Lookup tables PRODUCT_RATES = { 1: {"individual": 1000, "corporate": 900}, 2: {"individual": 1500, "corporate": 1200}, 3: {"individual": 2000, "corporate": 1700}, } CATEGORY_MULTIPLIER = {"electronics": 1.2, "groceries": 0.8, "luxury": 1.5} ## The calculation The base receivable is `volume * PRODUCT_RATES[productid][status]`. Three discount rates then apply, and they add together: - a spring rate of 0.05 when season is "spring", otherwise 0.0; - a volume rate that is 0.05 for volume >= 100 plus an additional 0.05 for volume >= 1000 (the two thresholds are independent and stack, so volume >= 1000 earns the full 0.10); - a strategic rate of 0.05 when `customer` is "strategic", otherwise 0.0. The receivable after discount is the base times (1 - sum of those three rates). Two loadings are computed from the receivable after discount: - a region loading of 0.10 * receivable_after_discount when `region` is "emea", else 0.0; - a payment loading of 0.025 * receivable_after_discount when `payment type` is "credit-card", else 0.0. The commission is (region_loading + payment_loading) * CATEGORY_MULTIPLIER[product category]. The subtotal is `receivable_after_discount` + `commission`. The `tax verified` column is True or False; withhold 0.03 * subtotal when the sale has not been tax-verified, and 0.0 when it has been verified. `total_receivables` = subtotal - withholding.</pre>
L4 · Natural rules (tier) <pre># Pricing & commission calculation spec Write `process_calculations(csv_path)` that reads the CSV and returns one dict per row with the key `total_receivables`. ## Input columns Each row has these columns (header names may have extra surrounding spaces (strip leading/trailing whitespace only, and keep inner spaces like in `payment` type); compare text in lower case; convert `volume` and `productid` to numbers; an empty cell means the stated default): - `productid` : "1", "2", or "3" - `status` : "individual" or "corporate" - `volume` : an integer - `season` : "spring", "summer", "winter", or "autumn" - `customer` : "strategic" or "non-strategic" - `region` : "emea" or "non-emea" - `payment type` : "ach/wire" or "credit-card" - `transaction type` : "domestic", "international", or "cross-border" - `product category` : "electronics", "groceries", or "luxury" - `tax verified` : "yes" or "no" (an empty cell counts as "no") ## The calculation The per-unit price depends on `productid` and `status`. For productid 1 it is \$1,000 when status is individual and \$900 when corporate; for productid 2 it is \$1,500 / \$1,200; for productid 3 it is \$2,000 / \$1,700. The base receivable is `volume` times that per-unit price. Then three discounts come off, and they stack (add together): - **5%* if the `season` is spring; - a volume discount based on how much was ordered: under 100 units gets no volume discount, 100 to 999 units gets 5%, and 1000 units or more gets 10%; - **5%* if the `customer` is strategic. Apply the combined discount to get the receivable after discount: Two surcharges are then figured from the receivable after discount: - **0%* if the `region` is EMEA; - **2.5%* if the `payment type` is credit-card. Commission is the sum of those two surcharge amounts, scaled by a factor that depends on `product category`: electronics x1.2, groceries x0.8, luxury x1.5. The subtotal is the receivable after discount plus the commission. If `tax verified` is "no" (or blank), withhold **3%* of the subtotal; if it is "yes", withhold nothing. `total_receivables` is the subtotal minus the withholding.</pre>	L8 · Tables + tier + True/False <pre># Pricing & commission calculation spec `process_calculations(csv_path)` reads a CSV and returns one dict per row, each with the key `total_receivables`. ## Input columns Each row has these columns (header names may have extra surrounding spaces (strip leading/trailing whitespace only, and keep inner spaces like in `payment` type); compare text in lower case; convert `volume` and `productid` to numbers; an empty cell means the stated default): - `productid` : "1", "2", or "3" - `status` : "individual" or "corporate" - `volume` : an integer - `season` : "spring", "summer", "winter", or "autumn" - `customer` : "strategic" or "non-strategic" - `region` : "emea" or "non-emea" - `payment type` : "ach/wire" or "credit-card" - `transaction type` : "domestic", "international", or "cross-border" - `product category` : "electronics", "groceries", or "luxury" - `tax verified` : True or False (an empty cell counts as False) ## Lookup tables PRODUCT_RATES = { 1: {"individual": 1000, "corporate": 900}, 2: {"individual": 1500, "corporate": 1200}, 3: {"individual": 2000, "corporate": 1700}, } CATEGORY_MULTIPLIER = {"electronics": 1.2, "groceries": 0.8, "luxury": 1.5} ## The calculation The base receivable is `volume * PRODUCT_RATES[productid][status]`. Three discount rates then apply, and they add together: - a spring rate of 0.05 when season is "spring", otherwise 0.0; - a volume rate based on how much was ordered: 0.0 under 100 units, 0.05 from 100 to 999 units, and 0.10 at 1000 units or more; - a strategic rate of 0.05 when `customer` is "strategic", otherwise 0.0. The receivable after discount is the base times (1 - sum of those three rates). Two loadings are computed from the receivable after discount: - a region loading of 0.10 * receivable_after_discount when `region` is "emea", else 0.0; - a payment loading of 0.025 * receivable_after_discount when `payment type` is "credit-card", else 0.0. The commission is (region_loading + payment_loading) * CATEGORY_MULTIPLIER[product category]. The subtotal is `receivable_after_discount` + `commission`. The `tax verified` column is True or False; withhold 0.03 * subtotal when the sale has not been tax-verified, and 0.0 when it has been verified. `total_receivables` = subtotal - withholding.</pre>

Walking the ladder

- **L1 — Prescriptive pseudocode** (mirrors v2). Literal dicts, additive Rule 4 with "no elif", yes/no boolean, bare decimals, step-by-step recipe. The anchor: **reachable, 12/12**.
- **L2 — Declarative spec**. Drops the imperative "Start $x = 0.0$... do NOT use elif" recipe; states *what* each rule computes in plain prose. Tables still handed over. **Reachable, 12/12** — the recipe was not load-bearing.
- **L3 — Business prose**. Removes the literal dictionaries and bare decimals; rates now appear as a $\$/\%$ narrative the model must turn into its own data structures. **Never reaches more than 4/12**. This is the cliff.
- **L4 — Natural rules**. Re-arms the volume tier as natural bands ("under 100 ... 100–999 ... 1000+"). Already past the cliff: **best 3/12**.
- **L5 — Original dialect**. Also re-arms the True/False boolean — the article's original phrasing. **Best 4/12**.
- **L6 — Tables + natural tier**. Control: the L4 tier phrasing, but with the literal tables kept. **Reachable, 12/12**.
- **L7 — Tables + True/False**. Control: the L5 boolean, tables kept. **Reachable, 12/12**.
- **L8 — Tables + tier + True/False**. Control: *both* traps at once, tables kept. Still **reachable, 12/12**.

Rung	Spec dialect	Pass-rate	Best of 10	Reachable
L1	Prescriptive pseudocode	3/10	12/12	✓
L2	Declarative spec	4/10	12/12	✓
L3	Business prose	0/10	4/12	✗
L4	Natural rules (tier)	0/10	3/12	✗
L5	Original dialect (boolean)	0/10	4/12	✗
L6	Tables + natural tier	2/10	12/12	✓
L7	Tables + True/False	3/10	12/12	✓
L8	Tables + tier + True/False	2/10	12/12	✓

Table 1 · Best-of-ten per rung. Pass-rate is how many of ten attempts hit a full 12/12; "reachable" means at least one did. Tables present (L1, L2, L6, L7, L8) $\Rightarrow$ reachable; tables removed (L3, L4, L5) $\Rightarrow$ not.

4 What the ladder shows — a summary

The break is the **L2 $\rightarrow$ L3 boundary**, and it is sharper than the original two-trap story. Dropping the step-by-step recipe (L1 $\rightarrow$ L2) costs nothing. Dropping the literal lookup tables (L2 $\rightarrow$ L3) is the cliff. And the failures are *not* the tier or boolean at all: at L3–L5 the model handles both correctly in prose. Inspecting *which rows pass* at the failing rungs, it is always exactly the four rows that are non-emea and ach/wire — the rows whose region and payment loadings are both zero, i.e. **commission = 0**. Every wrong row has a non-zero loading, and the error is precisely the missing commission: for instance an electronics/EMEA row scores a ratio of $0.8929 = 1/1.12$, exactly the effect of dropping $\text{commission} = (\text{region} + \text{payment loading}) \times 1.2$ from the subtotal.

The load-bearing scaffold is the literal lookup tables, not the phrasing of the famous traps. With the tables present, every dialect is reachable (L1, L2, L6, L7, L8 $\rightarrow$ 12/12); with the tables removed, none is (L3, L4, L5). The tables are necessary and nearly sufficient; the tier and boolean traps are neither necessary nor

sufficient — even both together (L8) still pass. Once the model must build its own data structures from prose, what collapses is the multi-step *commission* chain, not the two single-line constructs.

Takeaway for a spec author. You may write the ten-rule spec in plain declarative prose — but you must (a) hand over the lookup tables verbatim and (b) name every multi-step derived quantity (here, the commission) as its own explicit sub-step. Narrating a multi-step formula in one breath is the thing a 3B cannot reliably reconstruct.

Two methodology notes. First, *single greedy decoding is a knife-edge*: a one-line change to a spec's title flipped L1's deterministic output from 12/12 to 0/12, so a small model must be evaluated by sampling, not one greedy roll. Second, *the instruction wrapper can dominate the result*: an ambiguous "strip spaces from headers" made the model delete the space inside `payment type` and crash, and "never crash" induced `try/except` that swallowed the real bug — both were removed so the wrapper is precise and identical for every rung.

5 Can a repair loop rescue the failing rungs?

The repository's standard loop generates code, scores it, and feeds the failure back for another attempt. Its feedback is a **scalar total diff** ("row 3: got X, expected Y") — but the L3 failure is a *localizable structural omission*: the whole commission chain is dropped. A single wrong number cannot say *which* of ten steps is missing, so the model flails. We therefore built `--repair-rules`, a **white-box loop**: a reference computes every intermediate in the pipeline, and on each failing row the loop reports the *first* step that diverges — e.g. "row 0: correct through `payment_loading`, but `commission` = 0.00, should be 3206.25" — asking the model to expose those intermediates so they can be compared.

We ran it on the three failing rungs (three trajectories, up to six iterations) as an A/B on the repair temperature. The localization works — the model does fix the flagged step. At the default temperature (0.6) the score paths *oscillate* ([4, 0, 1, 0, 0, 9]): fixing one step regresses another, and nothing converges. **Dropping the repair temperature to 0.3** makes the edits conservative enough to stick.

Rung	Spec dialect	best-of-10 (no repair)	repair-rules @ temp 0.6	repair-rules @ temp 0.3
L3	Business prose	4/12	8/12 (0/3)	7/12 (0/3)
L4	Natural rules (tier)	3/12	4/12 (0/3)	8/12 (0/3)
L5	Original dialect (boolean)	4/12	9/12 (0/3)	12/12 (1/3, iter 2)

Table 2 · White-box repair on the failing rungs. Cell shows best score and, in parentheses, trajectories (of three) that reached a full 12/12. At temperature 0.3 the boolean rung **L5 converges to 12/12 in a single repair iteration** (path [4, 12]); L4 climbs to 8/12; L3 still resists.

Interpretation. White-box, per-step feedback is a *real* lever — enough to rescue the boolean rung and nearly the tier rung at a low temperature — but **not a guarantee** for a fully de-scaffolded spec. Rebuilding the entire commission chain from pure prose (L3) remains out of reach in this budget, and the extra burden of emitting a dozen intermediate values every turn itself destabilizes a 3B's generation (the residual zero-scores are crashes). The dependable fix for L3 is still the *spec* — put the table back, or run a spec-repair loop that rewrites the prose toward the L1/L2 dialect — rather than out-arguing the model with cleverer diffs.

6 Conclusion

Yes — a tiny AI on an 8 GB MacBook can write real, numerically-correct business logic, *provided the specification is written in its dialect*. Holding the model fixed and moving only the words, we located the boundary precisely: correctness survives plain declarative prose but collapses the moment the author stops handing over the data tables and expects the model to assemble a multi-step derived quantity from narrative. The two constructs long blamed for the failure turn out to be handled fine; the true fragility is structural composition. A white-box repair loop can push the frontier outward — rescuing the boolean rung to a perfect score at a conservative temperature — but the reliable, general lever remains the one this line of work started from: **engineer the spec, not the model**.

Model: `mlx-community/Qwen2.5-Coder-3B-Instruct-4bit` via `mlx-lm 0.29.1` (Python 3.9, Apple Silicon). Dataset: 12 rows, tolerance 0.5. Ladder metric: best-of-10 at temperature 0.4. Repair: iteration-1 greedy, subsequent iterations sampled. All specs, code, datasets, and result JSONs are in `experiments/sweet-spot/`; this report is generated from those artifacts by `docs/build_report.py`.